

CS061:

Machine Organization & Assembly Language

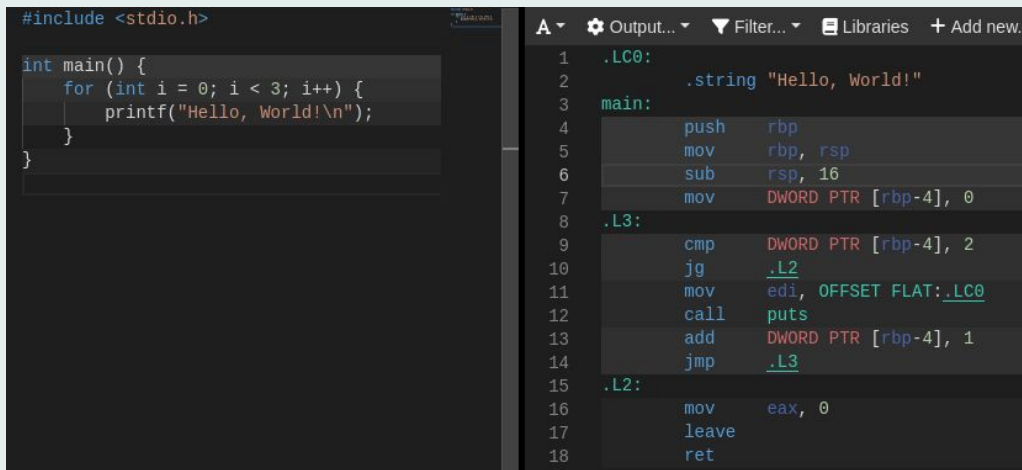
Lab 9

Agenda

1. Presentation:
 - a. x86-64 Crash Course
 - b. Compiler Optimizations
 - c. Out-of-order Execution
 - d. Godbolt Demo
 - e. Lab Descriptions
2. Work Time / Questions / Demos

Lab Overview

- Use Godbolt (an online compiler) to view the x86-64 assembly code of C & C++ programs.
- Explore how assembly code changes between:
 - Compiler Optimization Levels
 - C and C++
 - Loop Unrolling



The screenshot displays the Godbolt online compiler interface. On the left, the C source code is shown in a dark-themed editor. It includes a header file and a `main` function with a `for` loop that prints "Hello, World!" three times. On the right, the generated x86-64 assembly code is displayed, showing the compiler's internal labels and instructions for the same program.

```
#include <stdio.h>

int main() {
    for (int i = 0; i < 3; i++) {
        printf("Hello, World!\n");
    }
}
```

```
1  .LC0:
2      .string "Hello, World!"
3  main:
4      push    rbp
5      mov     rbp, rsp
6      sub     rsp, 16
7      mov     DWORD PTR [rbp-4], 0
8  .L3:
9      cmp     DWORD PTR [rbp-4], 2
10     jg      .L2
11     mov     edi, OFFSET FLAT:.LC0
12     call    puts
13     add     DWORD PTR [rbp-4], 1
14     jmp     .L3
15  .L2:
16     mov     eax, 0
17     leave
18     ret
```

x86-64 Crash Course

- x86 is a separate assembly language with a separate ISA
 - **ISA**: Instruction Set Architecture
 - **CISC**: Complex Instruction Set Computer
- x86 has 16 registers that hold 64 bits per register!
 - LC-3 has 8 registers that hold 16 bits per register.
- x86 has a lot of instructions (> 1000 instructions)
 - LC-3 only has 15 instructions.
- x86 instructions are not fixed-length!
 - LC-3 instructions are always 16-bits wide.
- x86 has a built-in stack!
 - Had to create our own stack in LC-3.

x86-64 Registers

- x86 has 16 registers each 64-bits wide
- Some x86 registers are named!
 - E.g. `%rax`, `%rcx`, `%rdx`, `%rbx`
 - E.g. `%r8`, `%r9`, `%r10`
- Can access portions of a register too!
 - `%rax` is the entire 8-byte (64-bit) register.
 - `%eax` access the first 4-bytes (0-3) of the `%rax` register!
 - `%ax` access the first 2 bytes of the `%rax` register!
 - `%al` access the first byte of the `%rax` register

		q (8 bytes)		1 (4 bytes)	w (2 bytes)	b (1 byte)	
<code>%rax</code>	<code>%eax</code>					<code>%ax</code>	accumulate
<code>%rbx</code>	<code>%ebx</code>					<code>%bx</code>	base
<code>%rcx</code>	<code>%ecx</code>					<code>%cx</code>	counter
<code>%rdx</code>	<code>%edx</code>					<code>%dx</code>	data
<code>%rsi</code>	<code>%esi</code>					<code>%si</code>	source index
<code>%rdi</code>	<code>%edi</code>					<code>%di</code>	destination index
<code>%rsp</code>	<code>%esp</code>					<code>%sp</code>	stack pointer
<code>%rbp</code>	<code>%ebp</code>					<code>%bp</code>	base pointer

x86 Instructions

- x86 has a lot of instructions (> 1000)!
- Example Instruction format:
 - Instruction <Destination Register> <Source Register>
 - Same as LC-3!
 - Not the only syntax: Intel vs AT&T
- Example instructions:
 - `add <dest reg> <src reg>`
 - Add the destination register with the source register.
 - `add rax, rcx` : `rax += rcx`
 - `mov <dest reg> <src reg>`
 - Move source register value to destination register.
 - `mov rbp, rsp` : `rbp = rsp`
 - `push rbp` : Push rbp register onto the stack

x86 Control Flow

- Control flags that are used by control flow instructions.
 - SF (Sign Flag): Set if result is negative (similar to n register in LC-3)
 - ZF (Zero Flag): Set if result is zero.
 - OF (Overflow Flag): Set if 2's complement overflow has occurred.
 - CF (Carry Flag): Unsigned overflow.
- Instructions that set flags:
 - `cmp <reg 1> <reg 2>`
 - Performs Reg 1 - Reg 2 and sets control flags.
 - `test <reg 1> <reg 2>`
 - Performs reg 1 & reg 2 and sets control flags.
 - Useful for testing if `reg 1 == reg 2`.

x86 Control Flow

- Lots of jump instructions!
- Examples:
 - `jmp label`: Jump to the label.
 - `je/jz label`: Jump if result is equal/zero (ZF = 1).
 - `jg/jnle label`: Jump if result is greater (signed >).
- Combined:
 - `test rax rbp`
 - `je label`
 - Jumps to label if `rax == rbp`.

Compiler Optimizations

- Optimization Levels:
 - How hard should the compiler try to optimize your code?
 - No optimization: -O0
 - Optimization Level 1 (-O1): Minor optimizations.
 - Levels 1 to 3
- Loop Unrolling:
 - Say we have a for loop that runs 3 times.
 - Each time it executes an add instruction.
 - Unrolling the loop means we just have 3 add instructions!
 - Why do this?

```
for (i = 0; i < 3; i ++):  
    x = x + 10
```

↓

```
x = x + 10  
x = x + 10  
x = x + 10
```

All is not as it seems!

EXTRA INFO - Not on the labs / exams!

- Myth: Code executes sequentially.
 - Sort-of!
 - Compilers can reorder your code!
 - CPU can execute instructions *out-of-order*!
- How can CPU handle if-statements and loops if instructions out of order?
 - CPU tries to predict what will happen! (*Branch Prediction*)
 - Right Prediction: Awesome!
 - Wrong Prediction: Costly!
 - Has to trash the instructions it tried to process out of order.

Exercise Overview

- **Exercise 1:**
 - Comparing optimization levels -O1 to -O3 for a C snippet.
- **Exercise 2:**
 - Comparing number of lines generated by C and C++.
 - Demangling Identifiers.
- **Exercise 3:**
 - Comparing normal generated code versus code with loops unrolled.

Demo Info

- **Lab Grade Breakdown:**
 - 3 points for attendance.
 - 7 points for demoing (+1 bonus point demo'd before/during Friday).
 - 3 point penalty if lab is demo'd during the next lab session.
- **Tips before you demo:**
 - ***Understand your code!*** (Know what each line does & the input/output)
 - ***Test your code!*** (Check for correct output and that there are no errors)